
Xiaoxiao (George) Li

(647) 830-3876

george.li.24601@gmail.com

[Github](#)

[Blog](#)

SKILLS

>3 years as the owner of projects serving > 25M users:

- Java - Microservice with Spring boot (> 1k TPS)
- Kafka - event sourcing
- AWS RDS Aurora - (> 10k qps with billions of records)
- TiDB (> 40 k qps)
- ElasticCache - cluster mode with HA

[See here for data points on soft skills](#)

EXPERIENCE

Paytm Labs/Paypay, Toronto, ON/Tokyo, Japan - *Software engineer*

June 2017 - Present

- As the system architect, I migrated payment service from Aurora to tidb, and increased the throughput by 3x.
 - This service is on the most critical path of > 25M users, processing hundreds of million dollars each day
 - [See my company's blog post](#)
- As the system architect, I migrated the transaction history service from Aurora to Tidb.
 - This service supports part of the app home screen (> 25 M users) and a first-level page with 5K API hits per sec
- Out of my own initiative, I reduced the infrastructure bill by \$150k per month combined. [See my blog post for my methodology](#)
- Firefought sister teams' SEV2s during multiple campaign days, each of which has > \$100M GMV
- Created and owned the reliable messaging framework that powers the event sourcing in multiple projects

MongoDB, New York City, NY - *Software engineer*

Jan 2017 - May 2017

- Billing and alerting system for MongoDB cloud

Paytm Labs, Toronto, ON - *Platform engineer*

Sep 2015 - Dec 2016

- Created and owned a Kafka-Samza pipeline that handles stream joins (1-1 and 1-n), stream enrichment, reprocessing, and out of order arrival.

-
- Incoming message rate at 20 mil/day.
 - This joined stream forms the base of other real time analytic streams. The results are projected 24/7 next to CEO's desk

Varicent, an IBM company, Toronto, ON - *Lead Software Developer*

Aug 2012 - Aug 2015

Varicent Software, Toronto, ON - *Software Developer*

May 2010 - Aug 2012

- Lead the migration of the main cash cow product from SQL Server to IBM Netezza.
 - Reduced the calculation time of the 3TB prod data set by $\frac{2}{3}$
 - This migration secured the company a deal > \$1M per year
- Building from scratch a domain specific process engine with Apache Kafka, Apache Samza, Scala. [Showed my work to Toronto Hadoop User Group](#)
- Won IBM Outstanding technical achievement award 2014

EDUCATION

University of Waterloo, Waterloo, ON - *Bachelor of Mathematics,*

2005 - 2010

Co-op, Computer Science with C & O Minor. Graduated with Distinction

Competitive Programming

- [1049 solves on leetcode as of Jan 17](#)
- Qualified for Google Code Jam round 2 2018
- [AtCoder max rating: 1611](#)
- [Codeforces max rating: 1802](#)
- [Topcoder max rating: 1411](#)